

Gradient descent

Phuong Luu Vo

December 14, 2022

Most of the optimization problems do not have a closed-form solution. Therefore, **iterative methods** are commonly used to find solutions to convex optimization problems.

In this lecture, we consider *first order algorithms* including gradient descent, stochastic gradient descent, subgradient, and proximal gradient descent to solve an unconstrained convex optimization problem.

1 Introduction

Consider the unconstrained optimization problem

$$\min_x f(x)$$

where f is convex and differentiable for all $x \in \text{dom}(f)$.

The necessary and sufficient condition for x^* be the optimal point is

$$\nabla f(x^*) = 0.$$

The iterative method generates a sequence $x^{(0)}, x^{(1)}, \dots, x^{(k)}, \dots \in \text{dom} f$ such that $\nabla f(x^*) \rightarrow 0$ when $k \rightarrow \infty$. Therefore, $f(x^{(k)}) \rightarrow p^*$, which is the optimal solution.

1.1 General descent method

The update

$$x^{(k+1)} = x^{(k)} + t_k \Delta x^{(k)}$$

such that $f(x^{(k+1)}) < f(x^{(k)})$ for any $k = 0, 1, \dots$ is called *descent method*. $\Delta x^{(k)}$ is *search direction* and t_k is *line search*.

From first-order condition,

$$f(y) \geq f(x) + \nabla f(x)^T (y - x).$$

Replacing $y = x^{(k+1)}$ and $x = x^{(k)}$ yields

$$f(x^{(k+1)}) \geq f(x^{(k)}) + t_k \nabla f(x^{(k)})^T \Delta x^{(k)}.$$

To guarantee $f(x^{(k+1)}) < f(x^{(k)})$, it must be

$$\nabla f(x^{(k)})^T \Delta x^{(k)} < 0. \quad (1)$$

Vector $\nabla f(x^k)$ and search direction Δx^k form an obtuse angle. Or search direction and $\nabla f(x)$ are in different halfspaces created by the hyperplane $\nabla f(x^k)^T(x - x^k) = 0$ (see Fig. 3).




Figure 1: *Search direction* in descent method. Δx and $\nabla f(x)$ are in different halfspace.

General descent method is as follows:

Algorithm 1 General descent method.

Initialize $x^{(0)} \in \text{dom} f$;

repeat

1. Determine *search direction* $\Delta x^{(k)}$;
2. *Line search*: choose step size $t_k > 0$.
3. Update: $x^{(k+1)} = x^{(k)} + t_k \Delta x^{(k)}$, $k = 1, 2, 3 \dots$

until stopping criterion satisfies.

Theoretically, stopping condition is $\|\nabla f(x)\|_2 \leq \epsilon$, in which ϵ is a very small number.

Choosing step size:

- *fix step size*: the simplest strategy
- *exact line search*: in every iteration, choose t such that

$$t = \operatorname{argmin}_{t > 0} f(x + t \Delta x).$$

This is not practical, and almost not used. However, in some special instances that the solution to the problem $\operatorname{argmin}_{t \geq 0} f(x + t\Delta x)$ is easily calculated, we can apply exact line search.

- *backtracking line search*: fix parameters $\alpha \in (0, 1/2)$, $\beta \in (0, 1)$ and $t = 1$, repeat $t := \beta t$ until $f(x + t\Delta x) \leq f(x) + \alpha t \nabla f(x)^T \Delta x$.

2 Gradient descent method (GD)

In *gradient descent* method (GD), search direction $\Delta x = -\nabla f(x)$. Therefore, gradient descent algorithm is as follows:

Algorithm 2 Gradient descent algorithm.

initialize $x^{(0)} \in \operatorname{dom} f$;

repeat

1. *Line search*: choose step size t_k .
2. Update: $x^{(k+1)} = x^{(k)} - t_k \nabla f(x^{(k)})$, $k = 1, 2, \dots$

until stopping criterion satisfies.

Theoretically, the stopping criterion is when $\|\nabla f(x)\|_2 \leq \epsilon$. The practical other criterion is $\|x^{(k)} - x^{(k-1)}\|_2 \leq \epsilon$.

2.1 Convergence analysis of GD

A function f is *Lipschitz continuous* if there is L such that

$$\|f(x) - f(y)\|_2 \leq L\|x - y\|_2. \quad (2)$$

If f is convex and differentiable with $\operatorname{dom}(f) = \mathbb{R}^n$, and additionally ∇f is L -Lipschitz continuous, GD converges with fixed step size $t \leq 1/L$,

$$f(x^{(k)}) - p^* \leq \frac{\|x^{(0)} - x^*\|_2^2}{2tk} = \frac{C}{k}.$$

The same result holds for backtracking line search with $t = \beta/L$.

If f is strongly convex, *i.e.*, $f(x) - \frac{m}{2}\|x\|_2^2$ is convex for some $m > 0$, GD with fix step size $t \leq 2/(m + L)$ or with backtracking line search satisfies

$$f(x^{(k)}) - p^* \leq c^k \frac{L}{2} \|x^{(0)} - p^*\|_2^2,$$

in which $0 < c < 1$ depending on $m, x^{(0)}$.

3 Gradient descent applied in some problems

3.1 Linear regression

The least-square problem in linear regression is given by

$$\min. L(w) = \sum_{i=1}^N (y_i - x_i^T w)^2 = \|y - Xw\|_2^2. \quad (3)$$

in which $X = \begin{bmatrix} x_1^T \\ \dots \\ x_N^T \end{bmatrix}$ and $y = \begin{bmatrix} y_1 \\ \dots \\ y_N \end{bmatrix}$ (N data points).

Gradient of $L(w)$ is

$$\nabla L(w) = 2 \sum_{i=1}^N x_i (x_i^T w - y_i) = 2(X^T X w - X^T y)$$

Gradient descent algorithm is as follows:

Algorithm 3 Gradient descent applied in least square.

initialize w ;

repeat

1. choose step size t .
2. update:

$$w := w - 2t \sum_{i=1}^N x_i (x_i^T w - y_i)$$

in matrix form

$$w := w - 2t(X^T X w - X^T y)$$

until stopping condition satisfies.

3.2 Logistic regression

Binary classification using logistic regression is as follows:

$$\max. J(w) = \sum_{i=1}^N J_i(w) = \sum_{i=1}^N y_i x_i^T w - \log(1 + e^{x_i^T w}),$$

We have gradient

$$\begin{aligned} J_i(w) &= y_i x_i - \frac{e^{x_i^T w}}{1 + e^{x_i^T w}} * x_i \\ &= (y_i - \sigma_i) x_i, \end{aligned}$$

in which $\sigma_i = \frac{e^{x_i^T w}}{1 + e^{x_i^T w}}$.

Gradient descent applied in logistic regression is as follows:

Algorithm 4 Gradient descent applied in logistic regression.

initialize w ;

repeat

1. choose step size t .

2. update:

$$w := w + t \sum_{i=1}^N (y_i - \sigma_i) x_i$$

until converge.

Some notes:

- In above example, the optimization is maximizing the log likelihood, hence, gradient update must be

$$w := w + t \nabla J(w). \quad (4)$$

- In many applications, especially in deep learning, gradient is calculated numerically. E.g., backward function in PyTorch can calculate the gradient if x is a tensor with `requires_grad = True`.

The following code implements GD in logistic regression from scratch with artificial dataset.

Import library:

```
import numpy as np
import matplotlib.pyplot as plt
```

Define sigmoid function:

```
def sigmoid(s):
    return 1/(1+ np.exp(-s))
```

Generate dataset with N data points, D parameters. X with size $N \times D$ follows normal distribution. Y is generated using logistic model with w_{true} plus some noise.

```
np.random.seed(1)
D = 10    # number of parameters
N = 100   # number of samples
w_true = np.random.randn(D,1)
X = np.random.normal(0, 5, size=(N,D))
Y = np.ROUND(sigmoid(X.dot(w_true) \
                    + np.random.normal(0,1, size=(N,1))))
```

Verify sizes of X, Y

```
print(X.shape)
print(Y.shape)
```

```
(100, 10)
(100, 1)
```

Gradient of J_i

```
def gradient_i(w, xi, yi):
    sigma_i = sigmoid(np.dot(xi, w))
    return xi*(yi - sigma_i)
```

Step size, number of iterations and initialize w

```
step_size = 0.005
iters = 100
```

```
w_init = np.random.rand(D, 1)
w = w_init.T
obj=np.array([])
w_last = w_init
count = 0
```

obj , w store the values of objective and w for all iterations to visualize the convergence of GD.

Gradient descent update:

```
while count < iters:
    # calculate gradient
    grad = np.zeros((D,1))
    for i in range(N):
        xi = X[i]
        yi = Y[i]
        grad += gradient_i(w_last, xi, yi).reshape(D,1)

    # perform gradient update and append to w
    w_cur = w_last + step_size * grad
    w = np.vstack((w, w_cur.T))

    # calculate objective value and append to obj
    obj_cur = np.sum(Y*np.dot(X, w_cur) - \
        np.log(1+np.exp(np.dot(X, w_cur))), axis=0)
    obj = np.append(obj, obj_cur)

    w_last = w_cur
    count += 1
```

We can also calculate the objective by the sum of J_i :

```

obj_cur = 0
for i in range(N):
    xi = X[i]
    yi = Y[i]
    obj_cur += yi*np.dot(xi.T, w_cur) - \
        np.log(1+np.exp(np.dot(xi.T, w_cur)))

```

Gradient J is calculated in matrix form:

```

sigma = sigmoid(np.dot(X,w))
grad = np.sum(X*(Y-sigma), axis=0)

```

Print the last value `print("w*",w_last)`.

Visualize:

```

plt.plot(obj, label="step_size: {}".format(step_size))
plt.xlabel("iterations")
plt.ylabel("log_likelihood")
plt.legend()
plt.savefig(' ./GD_obj_s{}.png'.format(step_size))

```




Figure 2: GD with $\text{step_size} = 0.01$.

References

- [1] Stephen P. Boyd and Lieven Vandenberghe. Convex optimization. Cambridge university press, 2004.
- [2] Convex optimization lectures. Available at www.stat.cmu.edu/~ryantibs/convexopt/.




Figure 3: GD with $\text{step_size} = 0.001$.